

Homework4

Name: Yang Cai

Program Discussion

Transformation of the data

Step 1 reads input file: `SparkContext.textFile: String => RDD[String]`

Step 2 parses each line to adjacent list:

`flatMap: RDD[String] => RDD[(String, java.util.List<String>)]`

Step 3 converts from Java list to Scala list and remove self-links:

`map: RDD[(String, java.util.List<String>)] => RDD[(String, List[String])]`

Step 4 merges all pairs that have same key, generate and persist this graph

`reduceByKey: RDD[(String, List[String])] => RDD[(String, List[String])]`

Step 5 count the number of nodes: `graph.count: RDD[(String, List[String])] => Int`

Step 6 generates the initial page rank value:

`mapValues: RDD[(String, List[String])] => RDD[(String, Double)]`

Loop starts:

Step 7 joins the graph and the pagerank:

`join: RDD[(String, List[String]), RDD[(String, Double)]] => RDD[(String, (List[String], Double))]`

Step 8 emits each page's contribution to its out-link, meanwhile add dangling nodes' rank to the global counter: `flatMap: RDD[(String, (List[String], Double))] => RDD[(String, Double)]`

Step 9 aggregates all contributions to the same page

`foldByKey: RDD[(String, Double)] => RDD[(String, Double)]`

Step 10 retrieves the dangling nodes' rank from global counter, use `RDD.count` to force an action (otherwise global counter won't be retrieved from driver program)

Step 11 calculate the final page rank value according to the formula and dangling nodes' rank.

`mapValues: RDD[(String, Double)] => RDD[(String, Double)]`

Loop Ends

Repeat Step 7 to Step 11 10 times

Step 12 take top 100 pages: `takeOrdered: RDD[(String, Double)] => Array[(String, Double)]`

Step 13 save this final output to text file: parallelize and `saveAsTextFile`:

`Array[(String, Double)] => RDD[(String, Double)] => save to file`

Narrow/Wide Steps:

Narrow steps are, Step 2, Step 3, Step 5, Step 6, Step 7, Step 8, Step 10, Step 11, Step 12

Wide steps are Step 4, Step 9

There are total 24 stages: 1 for Step 1(SparkContext. textFile), 1 for Step 4(reduceByKey), 20 for the 10-times loop(each loop has 2 stages, the i'th loop foldByKey, and also the i – 1' th loop's foldByKey is also counted), 1 when calling takeOrdered(this triggers a stage since it needs result of foldByKey of the 10th loop), and 1 for Step 13(parallelize)

Performance Comparison

1. Hadoop 5 workers: 4394 seconds
2. Hadoop 10 workers: 2762 seconds
3. Spark 5 workers: 2570 seconds
4. Spark 10 workers: 1228 seconds

The Spark is relative faster, it might because Spark does not read/write from/to HDFS again and again, so it has smaller I/O overhead and everything is in memory.